



APPLIED TIME SERIES ANALYSIS

Assignment 1

Group Number: 07

Group Member:

He Guanzhou	A0153091M
Hu Hao	A0152988L
Hu Jiongyi	A0153190M
Jiang Zeyu	A0152953B
Ye Rong	A0153626E

Exercise 1 (Electricity Forecast)

1. Load the data and use the command `ts(...)` to convert it into a time series object:

```
electricity_load <- read.table("E:/StatSoft/electricity_load.dat", quote="\"", comment.char="")  
head(electricity_load)
```

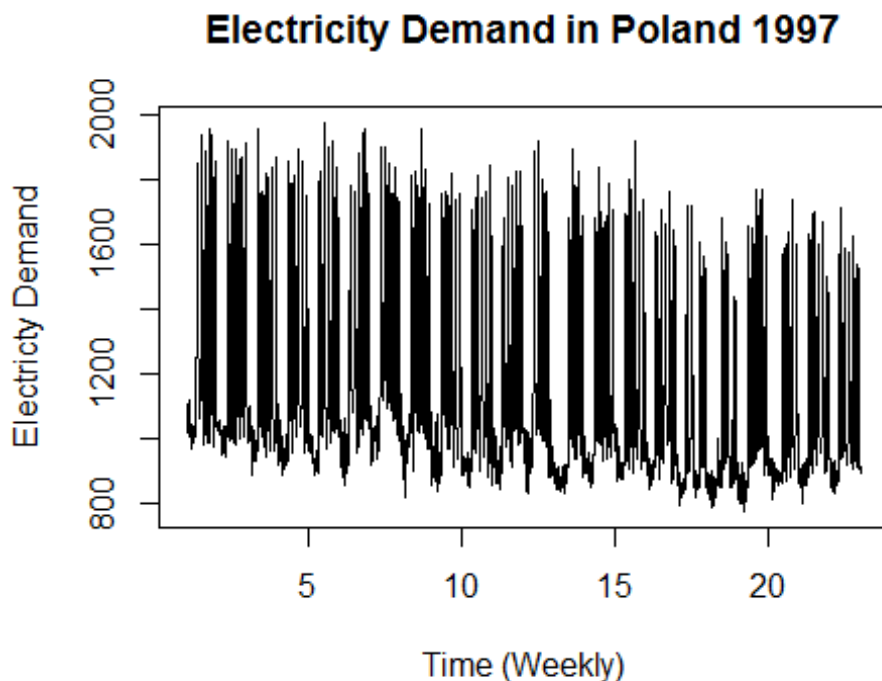
```
##      V1  
## 1 1087  
## 2 1032  
## 3 1018  
## 4 1031  
## 5 1057  
## 6 1024
```

```
electricity_ts<-ts(electricity_load$V1, start = 1, frequency = 24*7)  
head(electricity_ts)
```

```
## [1] 1087 1032 1018 1031 1057 1024
```

2. Display the time series:

```
plot(electricity_ts, ylab="Electricity Demand", xlab = "Time (Weekly)",  
main="Electricity Demand in Poland 1997")
```



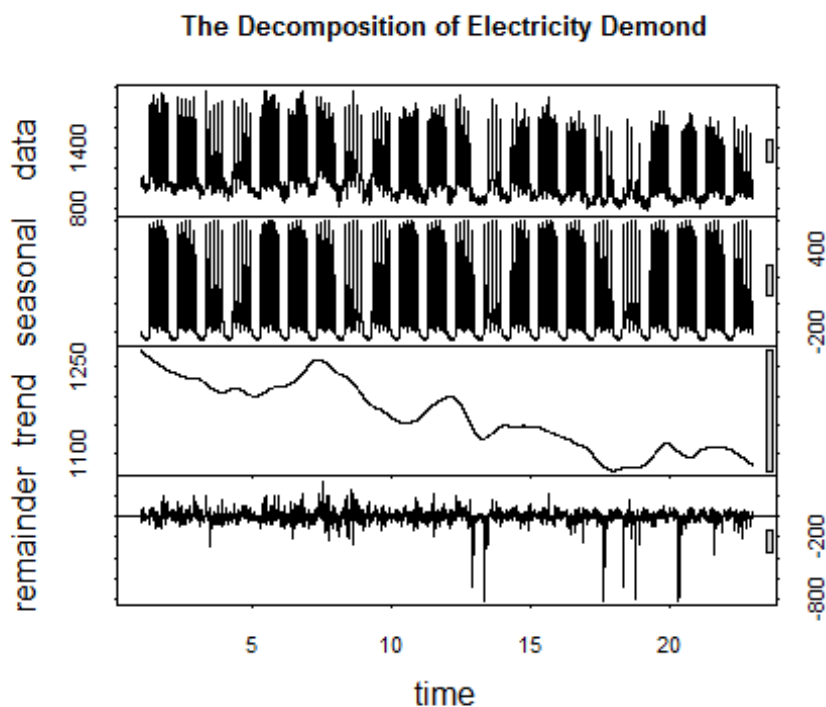
As we can see in the graph, an appropriate seasonal period is ONE unit on the Y-axis. Because we choose "24 * 7" as the frequency, the period should be ONE WEEK according to the data.

- Decompose the time series into the superposition of a trend / seasonal / remainder pattern:

```
## Loading required package: forecast
## Loading required package: zoo
##
## Attaching package: 'zoo'
##
## The following objects are masked from 'package:base':
##
##   as.Date, as.Date.numeric
## Loading required package: timeDate
## This is forecast 7.1
## Loading required package: fma
## Loading required package: tseries
## Loading required package: expsmooth
## Loading required package: lmtest
```

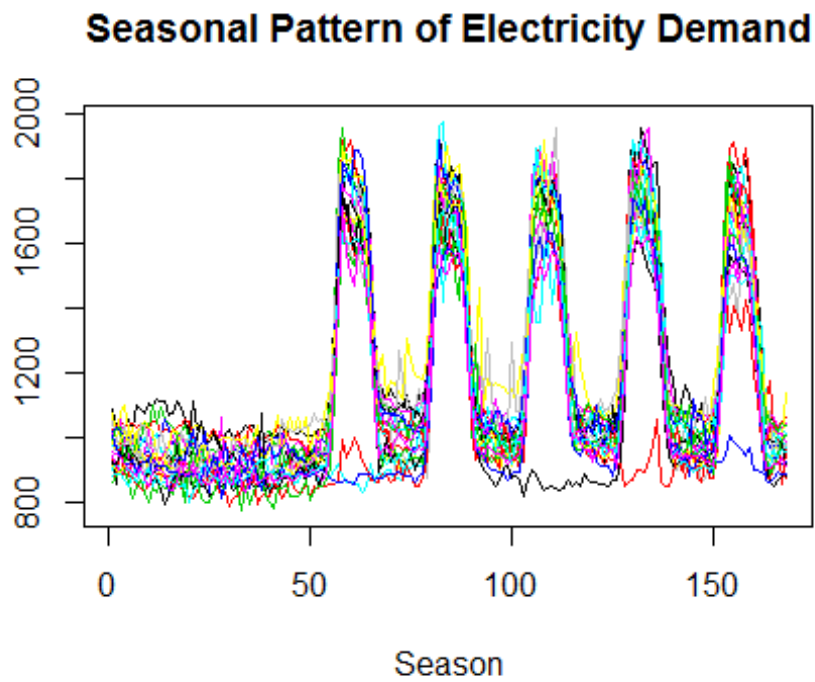
We use package "fpp" to decompose the time series.

```
fit<-stl(electricity_ts, s.window="periodic", robust=T)
plot(fit, main="The Decomposition of Electricity Demand")
```




4. Display the seasonal pattern:

```
library(forecast)
seasonplot(electricity_ts, s=24*7,col=1:23,type="l", main="Seasonal Pattern of Electricity Demand")
```



```
seasonplot(electricity_ts-fit$time.series[, "trend"], s=24*7,col=1:23,type="l", main="Seasonal Pattern (Remove Trend)")
```


Seasonal Pattern (Remove Trend)



In the remainder pattern, we can see there are some huge residual value at some point. Thus, the time series is not perfectly seasonal.

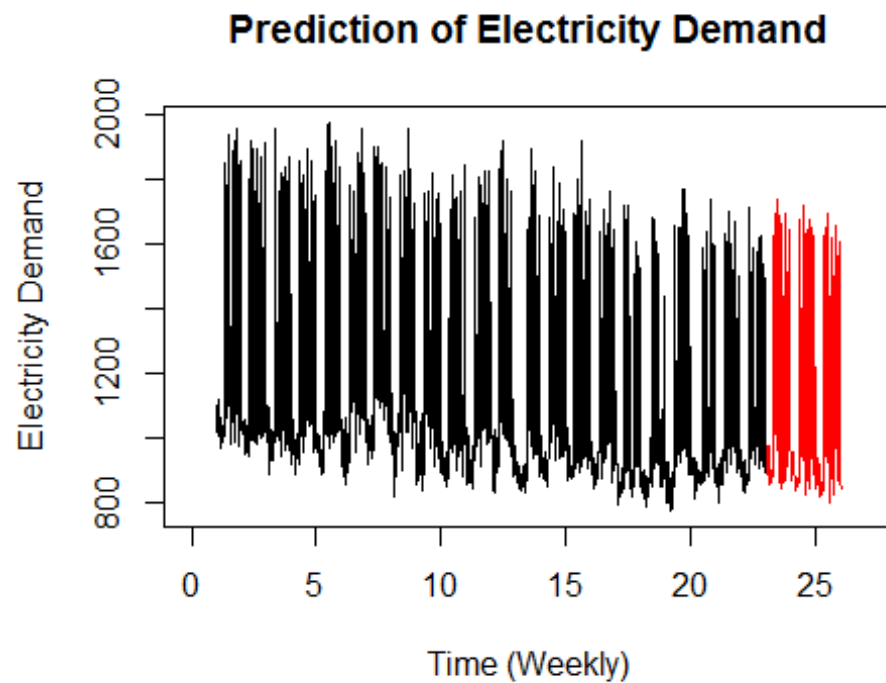
```
plot(fit$time.series[, "remainder"], main="Remainder Pattern of Electricity Demand", ylab="Remainder", xlab="Time (Weekly)")
```

Remainder Pattern of Electricity Demand



5. Make a forecast for the next 3 weeks:

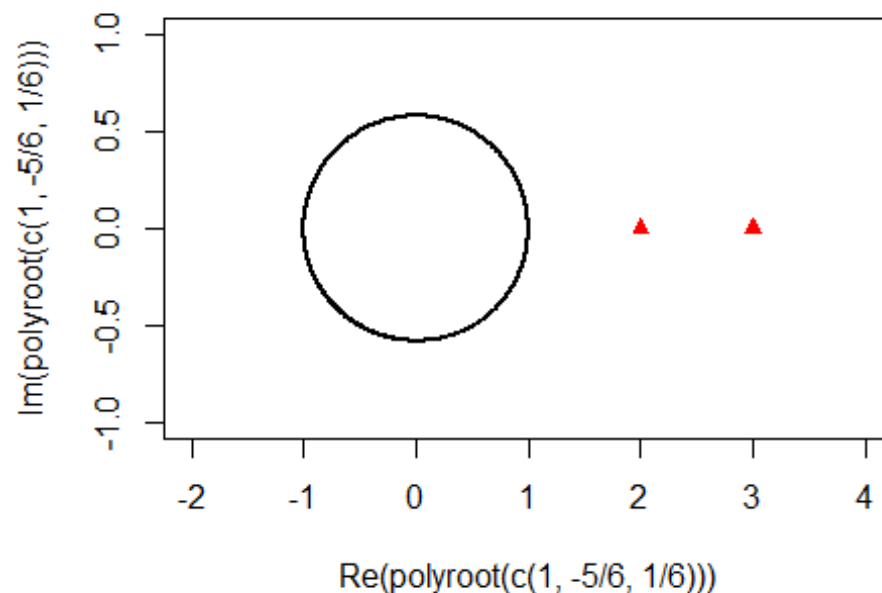
```
ele_optim<-hw(electricity_ts, initial = "simple", seasonal = "additive", h=24*7*3)
plot(electricity_ts, main="Prediction of Electricity Demand", ylab="Electricity Demand", xlab="Time (Weekly)", xlim=c(0,27))
lines(ele_optim$mean,col="red", type="l", lwd=1)
```



Exercise 2 (Bootstrap Estimate)

1. The MA(2) Process is invertible. This is because all the roots of the characteristic polynomial are greater than one.

```
library(shape)
plot(polyroot(c(1, -5/6, 1/6)), pch=17, col="red", xlim=c(-2, 4), ylim=c(-1, 1))
plotcircle(r=1, mid=c(0, 0), lwd=2)
```



2. Simulate a trajectory and double check:

```
MA2<-arima.sim(n=1000, list(ma=c(-5/6,1/6)), sd=1)
MA2_test<-arima(MA2, order=c(0,0,2))
summary(MA2_test)

##
## Call:
## arima(x = MA2, order = c(0, 0, 2))
##
## Coefficients:
##          ma1      ma2  intercept
##      -0.8588  0.1915    0.0071
## s.e.   0.0309  0.0305    0.0106
##
## sigma^2 estimated as 1.014:  log likelihood = -1426.35,  aic = 2860.
71
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE
## MASE
## Training set -9.175537e-05 1.007037 0.8061325 -82.80622 355.2077 0.4
182559
##              ACF1
## Training set -0.002422805
```

The coefficients of MA1 and MA2 are closed to $-5/6$ and $1/6$, so the estimation procedure provides a good estimate.


3. Estimate the accuracy of the estimation procedure for short time series:

```
beta<-c()
for (i in 1:1000)
{
  MA2_short_time<-arima.sim(n=40, list(ma=c(-5/6,1/6)), sd=1)
  MA2_st_test<-arima(MA2_short_time, order=c(0,0,2))
  beta[i]<-MA2_st_test$coef[2]
}
head(beta)

## [1] 0.240748317 0.183289360 0.144419602 0.578006280 -0.006454145
## [6] 0.102763945
```

4. Summary of the Estimates of Beta:

```
hist(beta, nclass=50, probability = T, col="steelblue", xlab="Beta", ma
in="Estimates of the Distribution of Beta")
```



```
mean(beta)

## [1] 0.1497037

var(beta)

## [1] 0.04911213
```

5. Estimate of Alpha:

```

MA1<-arima.sim(n=10000, list(ma=c(3)), sd=1)
MA1_test<-arima(MA1, order=c(0,0,1))
summary(MA1_test)

##
## Call:
## arima(x = MA1, order = c(0, 0, 1))
##
## Coefficients:
##          ma1  intercept
##          0.3312   -0.0061
## s.e.    0.0094    0.0400
##
## sigma^2 estimated as 9.013:  log likelihood = -25183.04,  aic = 5037
2.08
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MA
SE
## Training set 3.98034e-05 3.002244 2.397005 70.78541 202.8658 0.80733
23
##              ACF1
## Training set 0.001553553

```

That is because we use the Autocorrelation Function to estimate the coefficient Alpha, and the model with Alpha = 3 and Alpha=1/3 have the same ACF. However, only the model with Alpha = 1/3 is invertible. So the "arima" function in R shows that Alpha is closed to 1/3.

Exercise 3 (Sale Forecast)

1. Get the data and generate the time series: We download the data for Google Trend. The data starts in 2004, counting monthly.

```

aircon_data<-read.csv("E:/StatSoft/air_con.csv", header = FALSE, skip=
3)
head(aircon_data)

##          V1 V2
## 1 2004-01 13
## 2 2004-02 18
## 3 2004-03 16
## 4 2004-04 44
## 5 2004-05 34
## 6 2004-06 28


aircon_ts<-ts(aircon_data$V2, start = c(2004,1), frequency = 12)
head(aircon_ts)

## [1] 13 18 16 44 34 28


```

2. Plot the time series and decompose it:

```
plot(aircon_ts, main="Number of Queries Containing 'aircon'", ylab="Number", xlab="Time (Monthly)")
```




```
fit_aircon<-stl(aircon_ts, s.window = "periodic", robust=T)  
plot(fit_aircon, main="The Decomposition of Aircon")
```



As we can see in the graph, the time series has a trend and seasonal pattern. Therefore, we choose the Triple Exponential Smoothing method.

3. Fit and predict the time series:

```
aircon_optim<-hw(aircon_ts, initial = "optimal", seasonal = "additive",
h=20)
plot(aircon_ts, lwd=2, xlim=c(2011.6,2018), main="Prediction of the Air
con")
lines(fitted(aircon_optim), col="blue", type="l", lwd=2)
lines(aircon_optim$mean, col="red", type="l", lwd=2)
```



Then we find the number in June 2016, and the prediction in June 2017:

aircon_ts

##		Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
## 2004	13	18	16	44	34	28	25	13	22	25	24	22	
## 2005	15	23	18	21	32	31	27	29	30	19	26	24	
## 2006	17	22	25	27	24	32	30	28	24	34	21	24	
## 2007	24	23	27	31	30	35	26	24	21	28	20	21	
## 2008	21	19	19	31	41	31	21	21	20	23	20	21	
## 2009	18	19	22	30	31	44	26	28	25	30	21	21	
## 2010	18	30	32	34	43	34	30	27	29	34	28	26	
## 2011	24	28	29	34	52	38	40	35	37	33	30	32	
## 2012	32	34	36	42	52	56	40	37	40	42	34	32	
## 2013	37	33	51	55	58	100	58	44	44	43	37	38	
## 2014	34	42	51	58	67	83	67	48	52	53	46	46	
## 2015	43	39	52	67	74	81	69	56	62	64	53	54	
## 2016	61	55	70	93	84	68	76	63	58	71			

aircon_optim

##		Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## Nov 2016		64.16361	55.74601	72.58121	51.29000	77.03723
## Dec 2016		64.63915	56.17921	73.09909	51.70079	77.57751
## Jan 2017		63.41915	54.91170	71.92659	50.40813	76.43016
## Feb 2017		65.09782	56.53746	73.65819	52.00587	78.18977
## Mar 2017		71.10588	62.48697	79.72480	57.92440	84.28737

## Apr 2017	79.28684	70.60355	87.97014	66.00689	92.56679
## May 2017	84.47146	75.71776	93.22516	71.08384	97.85909
## Jun 2017	90.08184	81.25155	98.91213	76.57707	103.58660
## Jul 2017	78.31734	69.40412	87.23057	64.68574	91.94894
## Aug 2017	72.88235	63.87973	81.88498	59.11402	86.65068
## Sep 2017	74.50885	65.41024	83.60747	60.59372	88.42398
## Oct 2017	76.47316	67.27176	85.67457	62.40083	90.54550
## Nov 2017	71.10176	61.79093	80.41258	56.86208	85.34143
## Dec 2017	71.57729	62.15024	81.00435	57.15986	85.99473
## Jan 2018	70.35729	60.80716	79.90743	55.75163	84.96296
## Feb 2018	72.03597	62.35588	81.71605	57.23155	86.84038
## Mar 2018	78.04403	68.22711	87.86096	63.03034	93.05772
## Apr 2018	86.22499	76.26435	96.18563	70.99151	101.45847
## May 2018	91.40961	81.29840	101.52082	75.94585	106.87337
## Jun 2018	97.01998	86.75138	107.28859	81.31551	112.72446

So the number of queries in June 2016 is 68, and the prediction in June 2017 is 90.08, with the 95% prediction interval [76.58,103.59].


We know that the retailer sold 51 devices in June 2016, so the demand in June 2017 could be $51 / 68 * 90.08 = 67.56$, with the 95% prediction interval [57.435,77.69].

Exercise 4

1. Estimate the first three autocorrelation coefficients.

First of all, we load the data and generate the time series.

```
asg_1_AR3 <- read.table("E:/StatSoft/asg_1_MA3.dat", quote="", commen
t.char="")
AR3<-ts(asg_1_AR3$V1)
ACFAR3<-acf(AR3)
```



```
head(ACFAR3$acf)
## [1] 1.00000000 0.41980461 0.37193302 0.16221226 0.13283237 0.05947327
```

So the first three autocorrelation coefficients are 0.4198, 0.3719, 0.1622.

2. The associated Yule-Walker equations:

$$\rho(k) = \alpha\rho(k-1) + \beta\rho(k-2) + \gamma\rho(k-3)$$

3. Solve the Yule-Walker equations:

$$\begin{pmatrix} \rho(2) & \rho(1) & \rho(0) \\ \rho(1) & \rho(0) & \rho(1) \\ \rho(0) & \rho(1) & \rho(2) \end{pmatrix} \begin{pmatrix} \alpha \\ \beta \\ \gamma \end{pmatrix} = \begin{pmatrix} \rho(3) \\ \rho(2) \\ \rho(1) \end{pmatrix}$$

```
coeffi<-matrix(c(0.3719,0.4198,1,0.4198,1,0.4198,1,0.4198,0.3719),nrow
= 3)
rightside<-matrix(c(0.1622,0.3719,0.4198),nrow=3)
solution<-solve(coeffi)%*%rightside; solution
##           [,1]
## [1,] 0.33736780
## [2,] 0.26081680
## [3,] -0.07275798
```

Thus, $\hat{\alpha} = 0.34$, $\hat{\beta} = 0.26$, $\hat{\gamma} = -0.07$.